# Text Processing in Linux A Tutorial for NLP (CSE 562/662)

#### Kristy Hollingshead Winter 2007

www.cslu.ogi.edu/~hollingk/NLP\_tutorial.html

### Overview

- The goal here is to make your lives easier!
- NLP is very text-intensive
- Simple tools for text-manipulation
  - sed
  - awk
  - bash/tcsh
  - split
  - sort
  - head, tail
- When & how to use each of these tools

# Regular expressions crash course

- [a-z] exactly one lowercase letter
- [a-z]\* zero or more lowercase letters
- [a-z]+ one or more lowercase letters
- [a-zA-Z0-9] one lowercase or uppercase letter, or a digit
- [^(] match anything that is *not* '('

3

### sed: overview

- a stream editor
- WHEN
  - "search-and-replace"
  - great for using regular expressions to change something in the text
- HOW
  - sed 's/regexp/replacement/g'
    - 's/... = substitute
    - .../g' = global replace (otherwise will only replace first occurrence on a line!)

# sed: special characters

the start of a line...

except at the beginning of a character
set (e.g., [^a-z]), where it
complements the set

\$ the end of a line

• & the text that matched the regexp

We'll see all of these in examples...

5

# sed: (simple) examples

- eg.txt =
  - The cops saw the robber with the binoculars
- sed 's/robber/thief/g' eg.txt
  - The cops saw the thief with the binoculars
- sed 's/^/She said, "/g' eg.txt
  - She said, "The cops saw the robber with the binoculars
- sed 's/^/She said, "/g' eg.txt | sed 's/\$/"/g' She said, "The cops saw the robber with the binoculars"

# sed: examples from the homework!

- eg2.txt =
   (TOP (NP (DT The) (NNS cops)) (VP (VBD saw) (NP (DT the)
   (NN robber)) (PP (IN with) (NP (DT the) (NNS binoculars)))))
- "remove the syntactic labels" hint!: all of (and only) the syntactic labels start with '(' cat eg2.txt | sed 's/([^]\* //g' | sed 's/)//g' The cops saw the robber with the binoculars
- "now add explicit start & stop sentence symbols (<s> and </s>, respectively)"

```
cat eg2.txt | sed 's/([^ ]* //g' | sed 's/)//g' | sed 's/^/<s> /g' | sed 's/$/ <\/s>/g'
```

<s> The cops saw the robber with the binoculars </s>

7

# sed: (more complicated) example

- eg2.txt =
   (TOP (NP (DT The) (NNS cops)) (VP (VBD saw) (NP (DT the)
   (NN robber)) (PP (IN with) (NP (DT the) (NNS binoculars)))))
- "show just the POS-and-word pairs: e.g., (POS word)"

```
cat eg2.txt | sed 's/([^ ]* [^(]/~&/g' |
sed 's/[^)~]*~/ /g' |
sed 's/^ *//g' |
sed 's/))*/)/g'
```

(DT The) (NNS cops) (VBD saw) (DT the) (NN robber) (IN with) (DT the) (NNS binoculars)

#### awk: overview

- a simple programming language specifically designed for text processing
  - somewhat similar in nature to Tcl
- WHEN
  - using simple variables (counters, arrays, etc.)
  - treating each word in a line individually
- HOW

```
- awk 'BEGIN {initializations}
    /regexp1/ {actions1}
    /regexp2/ {actions2}
    END {final actions}' file.txt
```

(blue text indicates optional components)

9

# awk: useful constructions & examples

```
    each word in a line is a 'field'
    $1, $2, ..., $NF
    imagine every line of text as a row in a table; one word per column. $1 will be the word in the first column, $2 the next column, and so on up through $NF (the last word on the line)
```

- \$0 the entire row
- eq3.txt =

The cow jumped over the moon

- awk '{print \$2}' eg3.txtcow

The cow jumped over the 42 An old brown cow jumped over the 42

# awk: useful constructions & examples

- eg3.txt = The cow jumped over the moon
- if statements

```
- awk '{if ($1 == "he") { print $0; }}' eg3.txt
  (empty)
- awk '{if ($1 ~ "he") { print $0; } else { }}' eg3
```

- awk '{if (\$1 ~ "he") { print \$0; } else { ... }}' eg3.txt
  The cow jumped over the moon
  The
- for loops
  - awk '{for (j=1; j <= NF; j++) { print \$j }}' eg3.txt jumped</pre>
  - what if I only wanted to print every other word (each on a new line), in reverse order?
    - awk '{for (j=NF; j > 0; j=2) { print \$j }}' eg3.txt

1.1

12

cow

the

moon

# awk: useful constructions & examples

```
eq4.txt =
The cow jumped over the moon
                                                           1 The
And the dish ran away with the spoon
                                                           2 cow
printf statements
                                                           3 jumped
 - awk '{for (j=1; j <= NF; j++) { \
  printf("%d\t%s\n",j,$j);}}' eg4.txt</pre>
                                                           4 over
                                                           5 the
 – what if I want continuous numbering?
                                                           6 moon
 awk 'BEGIN {idx=0;} {for (j=1; j <= NF; j++) { \</pre>
    printf("%d\t%s\n",idx,\$j); idx++;}}' eg4.txt (1)And
                                                           2 the
substrings
 - substr(<string>, <start>, <end>)
   awk '{for (j=1; j \le NF; j+=2) {
    printf("%s ", substr($j,1,3))}; print "";}' eg4.txt
    The jum the
    And dis awa the
```

# awk: doing sed tasks with awk

- eg2.txt =
   (TOP (NP (DT The) (NNS cops)) (VP (VBD saw) (NP (DT the)
   (NN robber)) (PP (IN with) (NP (DT the) (NNS binoculars)))))
- "show just the POS-and-word pairs: e.g., (POS word)"
  cat eg2.txt | awk '{for (j=1;j<=NF;j++) {
   # if \$j is a word, print it (without its trailing paren's)
   if (substr(\$j,1,1) != "(") {
   i=index(\$j,")"); printf("%s ",substr(\$j,1,i))}
   # if \$j is a POS label, print it
   else {if (j+1<=NF &&
   substr(\$(j+1),1,1) != "(") printf("%s ",\$j)}}
   print ""}'
   (DT The) (NNS cops) (VBD saw) (DT the) (NN robber)
   (IN with) (DT the) (NNS binoculars)</pre>

13

### bash: overview

- shell script
- WHEN
  - repetitively applying the same commands to many different files
  - automate common tasks
- HOW
  - on the command line
  - in a file (type `which bash' to find your location):
     #!/usr/bin/bash
     <commands...>

# bash: examples

```
for f in *.txt; do
    echo $f;
    tail -1 $f >> txt.tails;
    done

for (( j=0; j < 4; j++ )); do
    cat part$j.txt >> parts0-3.txt;
    done

for f in hw1.*; do
    mv $f ${f//hw1/hw2};
    done
```

15

# miscellaneous

- sort
  - sort -u file.txt
    for a uniquely-sorted list of each line in the file
- split
  - cat file.txt | split -1 20 -d fold divide file.txt into files of 20 lines apiece, using "fold" as the prefix and with numeric suffixes
- WC
  - a counting utility
  - wc -[l|c|w] file.txt
    counts number of lines, characters, or words in a file

### miscellaneous

- head, tail
  - viewing a small subset of a file
  - head -42 file.txt
    for the first 42 lines of file.txt
  - tail -42 file.txt
    for the last 42 lines of file.txt
  - tail +42 file.txt
    for everything except the first 42 lines of file.txt
  - head -42 file.txt | tail -1
    to see the 42nd line of file.txt
- tr
  - "translation" utility
  - cat mixed.txt | tr [a-z] [A-Z] > upper.txt

17

# Putting it all together!

- Let's say I have a text file, and I'd like to break it up into 4 equally-sized (by number of lines) files.
- wc -l orig.txt
   8000
- the easy way:

```
cat orig.txt | split -d -l 2000 -a 1 - part;
for f in part*; do mv $f $f.txt; done
```

the hard way:

```
head -2000 orig.txt > part0.txt
tail +2001 orig.txt | head -2000 > part1.txt
tail +4001 orig.txt | head -2000 > part2.txt
tail -2000 orig.txt > part3.txt
```

## Putting it all together!

 Now for each of those files, I'd like to see a numbered list of all the capitalized words that occurred in each file... but I want the words all in lowercase.

```
for f in part*;
  do echo $f;
  cat $f | awk 'BEGIN {idx=0} {
    for (j=1; j <= NF; j++)
      if (substr($j,1,1) ~ "[A-Z]") {
        printf("%d\t%s\n", idx, $j);
        idx++;
      }
    }' - | tr [A-Z] [a-z] >
    ${f//part/out};
    echo ${f//part/out};
    done
```

19

# Putting it all together!

- Now I'd like to see that same list, but only see each word once (unique).
- hint: you can tell 'sort' which fields to sort on
- e.g., sort +3 -4 will skip the first 3 fields and stop the sort at the end of field 4; this will then sort on the 4<sup>th</sup> field. sort -k 4, 4 will do the same thing

```
for f in out*; do
  cat $f | sort +1 -2 -u > ${f//out/unique};
done
```

and if I wanted to re-number the unique lists?

```
for f in out*; do
  cat $f | sort -k 2,2 -u | awk 'BEGIN {idx=0}
  {$1=idx; print $0; idx++}' > ${f//out/unique};
done
```

# Putting it all together!

 And finally, I'd like to see the first 5 & the last 5 words in each list, but I already have a list of these first-and-lasts started, so I just want to add onto it instead of creating a new one.

```
for f in unique*; do
  head -5 $f >> top-and-bottom-5;
  tail -5 $f >> top-and-bottom-5; done
```

 (and of course, I could then re-number top-and-bottom-5 if I were so inclined)

21

#### Resources

- You can always look at the man page for help on any of these tools!
  - i.e.: `man sed', or `man tail'
- My favorite online resources:
  - sed: www.grymoire.com/Unix/Sed.html
  - awk: <u>www.vectorsite.net/tsawk.html</u>
  - bash: <u>www.tldp.org/LDP/abs/html/</u> (particularly section 9.2 on string manipulation)
- Google it. ☺

# Warning!

- These tools are meant for very simple textprocessing applications!
- Don't abuse them by trying to implement computationally-intensive programs with them
  - like Viterbi search and chart parsing
- Use a more suitable language like C, C++, or Java
  - another tutorial, on data structures for NLP, in February
  - start thinking about parse trees & nodes as classes
  - brush up on hash tables